

## Q1 Introduction to iptables

0 Points

# Introduction

*Firewall Exploration Lab Copyright © 2006 - 2021 by Wenliang Du. This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. If you remix, transform, or build upon the material, this copyright notice must be left intact, or reproduced in a way that is reasonable to the medium in which the work is being re-published.*

We will only be using **tasks 2-4** from the SEED lab, with additional **Task A, B, and C** not in the PDF file. The full lab detail can be found here: [Firewall Exploration Lab](#)

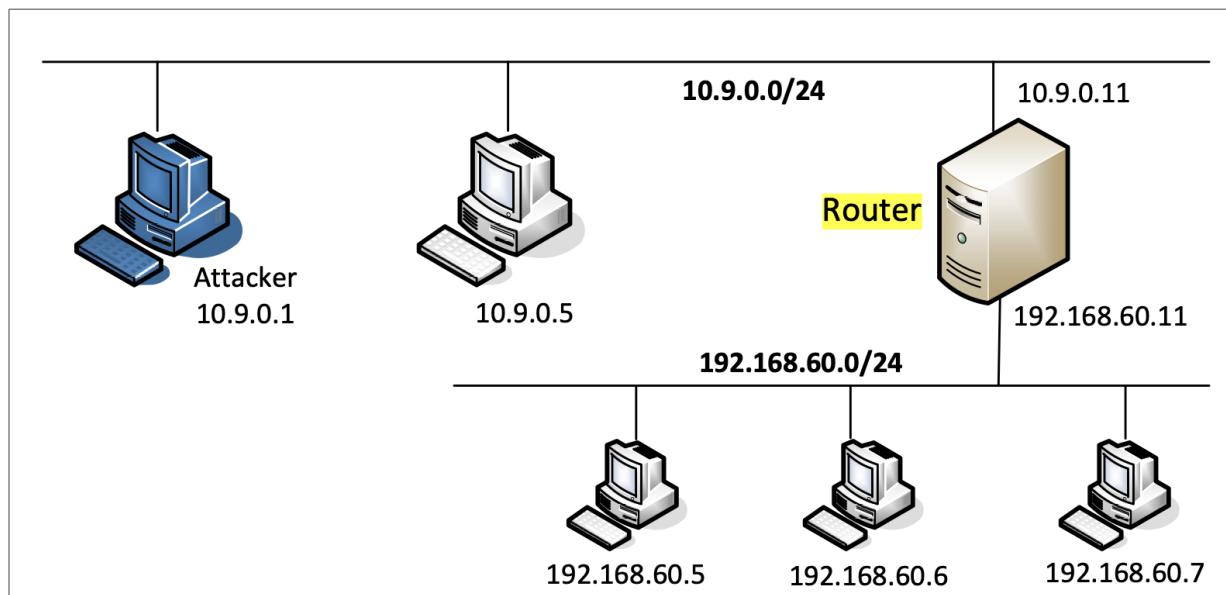
Earliest acceptance date: 30 April (+5% bonus)

Normal due date: 5 May

Latest acceptance date: 9 May (-4% points)

**Note:** Ensure all screenshots contain the full VM with the date and time.

Please follow the instructions in the [Firewall Explorations Lab](#) to setup the lab environment. Specifically, you would need to download the [Labsetup.zip](#) or [Labsetup-arm.zip](#) (for Apple Silicon) file and import it into the VM to set up the docker containers.



**Note: Please skip Task 1 and move to Task 2 (page 8)**

## **Q2 Task 2: Experimenting with Stateless Firewall Rules**

**35 Points**

Linux already has a built-in firewall, also based on netfilter. This firewall is called iptables. Technically, the kernel part implementation of the firewall is called Xtables, while iptables is a user-space program to configure the firewall. However, iptables is often used to refer to both the kernel-part implementation and the user-space program.

### **Using IPTables**

To add rules to the chains in each table, we use the iptables command, which is a quite powerful command. Students can find the manual of iptables by typing "man iptables" or easily find many tutorials from online. What makes iptables complicated is the many command-line arguments that we need to provide when using the command. However, if we understand the structure of these command-line arguments, we will find out that the command is not that complicated.

In a typical iptables command, we add a rule to or remove a rule from one of the chains in one of the tables, so we need to specify a table name (the default is filter), a chain name, and an operation on the chain. After that, we specify the rule, which is basically a pattern that will be matched with each of the packets passing through. If there is a match, an action will be performed on this packet. The general structure of the command is depicted in the following:

Table 1: `iptables` Tables and Chains

Table	Chain	Functionality
filter	INPUT FORWARD OUTPUT	Packet filtering
nat	PREROUTING INPUT OUTPUT POSTROUTING	Modifying source or destination network addresses
mangle	PREROUTING INPUT FORWARD OUTPUT POSTROUTING	Packet content modification

```
iptables -t <table> -<operation> <chain> <rule> -j <target>
```

-----  
**Table** Chain Rule Action

The rule is the most complicated part of the `iptables` command. We will provide additional information later when we use specific rules. In the following, we list some commonly used commands:

```
// List all the rules in a table (without line number)
iptables -t nat -L -n
// List all the rules in a table (with line number)
iptables -t filter -L -n --line-numbers
// Delete rule No. 2 in the INPUT chain of the filter table
iptables -t filter -D INPUT 2
// Drop all the incoming packets that satisfy the <rule>
iptables -t filter -A INPUT <rule> -j DROP
```

Note. Docker relies on `iptables` to manage the networks it creates, so it adds many rules to the `nat` table. When we manipulate `iptables` rules, we should be careful not to remove Docker rules. For example, it will be quite dangerous to run the "`iptables -t nat -F`" command, because it removes all the rules in the `nat` table,

including many of the Docker rules. That will cause trouble to Docker containers. Doing this for the filter table is fine, because Docker does not touch this table.

## Q2.1 Task 2.A: Protecting the Router

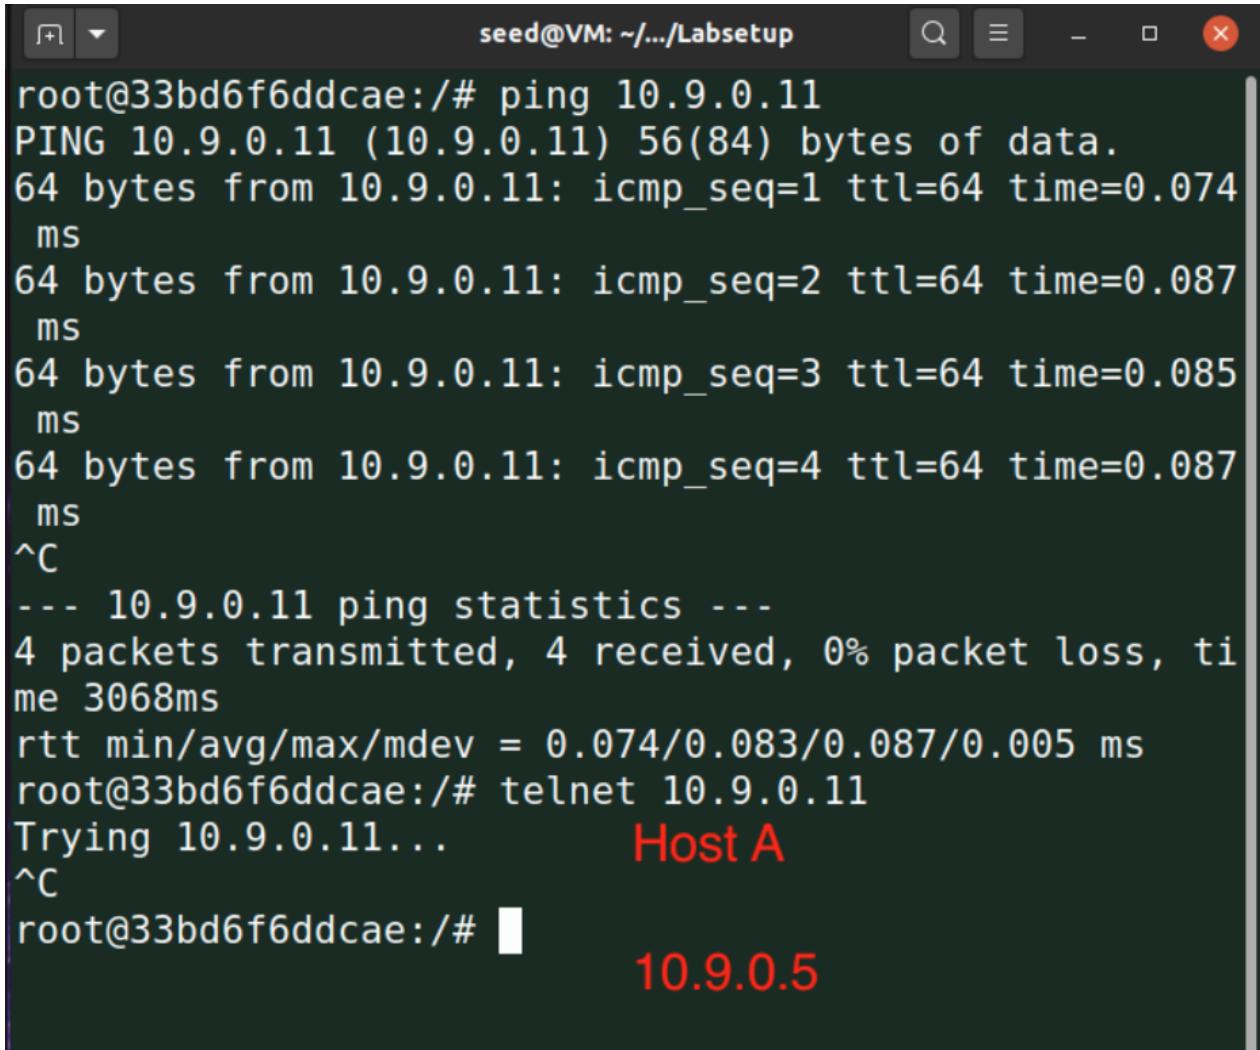
5 Points

In this task, we will set up rules to prevent outside machines from accessing the router machine, except ping. Please execute the following iptables command on the router container, and then try to access it from 10.9.0.5.

```
iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT  
iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT  
iptables -P OUTPUT DROP #Set default rule for OUTPUT  
iptables -P INPUT DROP #Set default rule for INPUT
```

### Expected Output

Note: Only ping packets are allowed, not telnet



The screenshot shows a terminal window with the following output:

```
seed@VM: ~/.../Labsetup  
root@33bd6f6ddcae:/# ping 10.9.0.11  
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.  
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.074 ms  
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.087 ms  
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.085 ms  
64 bytes from 10.9.0.11: icmp_seq=4 ttl=64 time=0.087 ms  
^C  
--- 10.9.0.11 ping statistics ---  
4 packets transmitted, 4 received, 0% packet loss, time 3068ms  
rtt min/avg/max/mdev = 0.074/0.083/0.087/0.005 ms  
root@33bd6f6ddcae:/# telnet 10.9.0.11  
Trying 10.9.0.11... Host A  
^C  
root@33bd6f6ddcae:/# 10.9.0.5
```

Upload screenshots below (must show ping and telnet)

```
[05/02/24] seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcdc host1-192.168.60.5
[05/02/24] seed@VM:~/.../Labsetup$ docksh 9a
root@9aa7af02f8b3:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.176 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.077 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.113 ms
64 bytes from 10.9.0.11: icmp_seq=4 ttl=64 time=0.149 ms
^C
--- 10.9.0.11 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3076ms
rtt min/avg/max/mdev = 0.077/0.128/0.176/0.037 ms
root@9aa7af02f8b3:/# telnet 10.9.0.11
Trying 10.9.0.11...
^C
root@9aa7af02f8b3:/#
```

### What is the purpose of each of the four rules?

- 1. If an incoming packet is an ICMP echo-request, let it pass through the firewall.
- 2. If an outgoing packet is an ICMP echo-reply, let it pass through the firewall.
- 3. If an outgoing packet does not match one of the rules, drop it.
- 4. If an incoming packet does not match one of the rules, drop it.

### Answer the following questions

- (1) Can you ping the router?
- (2) Can you telnet into the router?

- 1. Yes
- 2. No

**Cleanup.** Before moving on to the next task, please restore the filter table to its original state by running the following commands:

```
iptables -F  
iptables -P OUTPUT ACCEPT  
iptables -P INPUT ACCEPT
```

Another way to restore the states of all the tables is to restart the container. You can do it using the following command (you need to find the container's ID first):

```
$ docker restart <Container ID>
```

## **Q2.2 Task 2.B: Protecting the Internal Network**

**15 Points**

In this task, we will set up firewall rules on the router to protect the internal network 192.168.60.0/24. We need to use the FORWARD chain for this purpose.

The directions of packets in the INPUT and OUTPUT chains are clear: packets are either coming into (for INPUT) or going out (for OUTPUT). This is not true for the FORWARD chain, because it is bi-directional: packets going into the internal network or going out to the external network all go through this chain. To specify the direction, we can add the interface options using "-i xyz" (coming in from the xyz interface) and/or "-o xyz" (going out from the xyz interface). The interfaces for the internal and external networks are different. You can find out the interface names via the "ip addr" command.

In this task, we want to implement a firewall to protect the internal network. More specifically, we need to enforce the following restrictions on the ICMP traffic:

1. Outside hosts cannot ping internal hosts.
2. Outside hosts can ping the router.
3. Internal hosts can ping outside hosts.
4. All other packets between the internal and external networks should be blocked

**Note: "Outside" refers to the internet, not just the 10.9.0.11/24 network.**

You will need to use the "-p icmp" options to specify the match options related to the ICMP protocol. You can run "iptables -p icmp -h" to find out all the ICMP match options. The following example drops the ICMP echo request.

```
iptables -A FORWARD -p icmp --icmp-type echo-request -j DROP
```

**Expected Output** (Note: Ignore the 2.4 in the text)

Host A

```
root@772daac02cc:/# ifconfig
eth0: flags=4163<UP,BROADCAST,MULTICAST> mtu 1500
      link-layer ...
      ether 02:47:08:09:00:05 txqueuelen 0 (Ethernet)
      RX packets 1820 bytes 44426 (44.4 kB)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 2686 bytes 229972 (229.9 kB)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
      link-layer ...
      ether 00:00:00:00:00:00 txqueuelen 0 (Local Loopback)
      RX packets 12 bytes 1054 (1.0 kB)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 12 bytes 1054 (1.0 kB)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

**ifconfig**

Host 3

```
root@68aca2a0a7c4:/# ifconfig
eth0: flags=4163<UP,BROADCAST,MULTICAST> mtu 1500
      link-layer ...
      ether 02:47:08:09:00:07 txqueuelen 0 (Ethernet)
      RX packets 728 bytes 55266 (56.2 kB)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 1026 bytes 88401 (88.4 kB)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
      link-layer ...
      ether 00:00:00:00:00:00 txqueuelen 0 (Local Loopback)
      RX packets 1026 bytes 88401 (88.4 kB)
      RX errors 0 dropped 0 overruns 0 frame 0
      TX packets 30 bytes 1894 (1.8 kB)
      TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

**ifconfig**

seed@VM: ~

## Upload a screenshot of your iptable rules from the router

Run the `iptables -L -v` command to view iptable rules

<b>▼ 2B_rules.png</b>							<b>Download</b>
<b>root@bf0a51fb96a8:/# iptables -L -v</b>							
<b>Chain INPUT (policy DROP 0 packets, 0 bytes)</b>							
pkts bytes target prot opt in out source destination							
0 0 ACCEPT icmp -- any any anywhere anywhere							
icmp echo-request							
<b>Chain FORWARD (policy DROP 0 packets, 0 bytes)</b>							
pkts bytes target prot opt in out source destination							
0 0 DROP icmp -- eth0 eth1 anywhere anywhere							
icmp echo-request							
0 0 DROP icmp -- eth1 eth0 anywhere anywhere							
icmp echo-reply							
0 0 ACCEPT icmp -- eth1 eth0 anywhere anywhere							
icmp echo-request							
0 0 ACCEPT icmp -- eth0 eth1 anywhere anywhere							
icmp echo-reply							
<b>Chain OUTPUT (policy DROP 0 packets, 0 bytes)</b>							
pkts bytes target prot opt in out source destination							
0 0 ACCEPT icmp -- any any anywhere anywhere							
icmp echo-reply							

## Upload screenshots showing proof of checking rules 1-4 above (Please label each items)

▼ internal\_ping\_outside.png

 Download

```
[05/02/24] seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcfdc host1-192.168.60.5
[05/02/24] seed@VM:~/.../Labsetup$ docksh 26
root@267760e3bc58:/# ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=63 time=0.138 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=63 time=0.184 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=63 time=0.219 ms
64 bytes from 10.9.0.5: icmp_seq=4 ttl=63 time=0.071 ms
^C
--- 10.9.0.5 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3059ms
rtt min/avg/max/mdev = 0.071/0.153/0.219/0.055 ms
root@267760e3bc58:/# █
```

▼ internal\_telnet\_outside.png

 Download

```
[05/02/24] seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcfdc host1-192.168.60.5
[05/02/24] seed@VM:~/.../Labsetup$ docksh 26
root@267760e3bc58:/# telnet 10.9.0.5
Trying 10.9.0.5...
█
```

▼ outside\_ping\_internal.png

 Download

```
[05/02/24] seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcfdc host1-192.168.60.5
[05/02/24] seed@VM:~/.../Labsetup$ docksh 9a
root@9aa7af02f8b3:/# ping 192.168.60.6
PING 192.168.60.6 (192.168.60.6) 56(84) bytes of data.
^C
--- 192.168.60.6 ping statistics ---
17 packets transmitted, 0 received, 100% packet loss, time 16393ms
root@9aa7af02f8b3:/# █
```

▼ outside\_ping\_router.png

 Download

```
[05/02/24]seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcde host1-192.168.60.5
[05/02/24]seed@VM:~/.../Labsetup$ docksh 9a
root@9aa7af02f8b3:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.144 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.084 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.057 ms
64 bytes from 10.9.0.11: icmp_seq=4 ttl=64 time=0.101 ms
^C
--- 10.9.0.11 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3049ms
rtt min/avg/max/mdev = 0.057/0.096/0.144/0.031 ms
root@9aa7af02f8b3:/# █
```

▼ outside\_telnet\_internal.png

 Download

```
[05/02/24]seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcde host1-192.168.60.5
[05/02/24]seed@VM:~/.../Labsetup$ docksh 9a
root@9aa7af02f8b3:/# telnet 192.168.60.6
Trying 192.168.60.6...
█
```

## Q2.3 Task 2.C: Protecting Internal Servers

15 Points

In this task, we want to protect the TCP servers inside the internal network (192.168.60.0/24). More specifically, we would like to achieve the following objectives.

1. All the internal hosts run a telnet server (listening to port 23). Outside hosts can only access the telnet server on 192.168.60.5, not the other internal hosts.
2. Outside hosts cannot access other internal servers.
3. Internal hosts can access all the internal servers.
4. Internal hosts cannot access external servers.
5. In this task, the connection tracking mechanism is not allowed. It will be used in a later task

**Note: "Outside" refers to the internet, not just the 10.9.0.11/24 network.**

You will need to use the "-p tcp" options to specify the match options related to the TCP protocol. You can run "iptables -p tcp -h" to find out all the TCP match options. The following example allows the TCP packets coming from the interface eth0 if their source port is 5000.

```
iptables -A FORWARD -i eth0 -p tcp --sport 5000 -j ACCEPT
```

## Expected Output

The screenshot shows three terminal windows labeled "Host A", "Host 3", and "Host 1".

- Host A:** Shows the configuration of the "eth0" interface with IP 192.168.60.5 and subnet mask 255.255.255.0. It lists various statistics for traffic on this interface.
- Host 3:** Shows the configuration of the "eth0" interface with IP 192.168.60.5 and subnet mask 255.255.255.0. It lists various statistics for traffic on this interface.
- Host 1:** Shows the configuration of the "eth0" interface with IP 192.168.60.6 and subnet mask 255.255.255.0. It lists various statistics for traffic on this interface.

**Question 1:**  
All the internal hosts run a telnet server (listening to port 23). Outside hosts can only access the telnet server on 192.168.60.5, not the other internal hosts.

**Question 2:**  
Outside hosts cannot access other internal servers.

**Question 3:**  
Internal hosts can access all the internal servers.

**Question 4:**  
Internal hosts cannot access external servers.

## Upload a screenshot of your iptable rules from the router

Run the `iptables -L -v` command to view iptable rules

▼ 2C_rules.png							 Download	
root@bf0a51fb96a8:/# iptables -L -v								
Chain INPUT (policy DROP 0 packets, 0 bytes)								
pkts	bytes	target	prot	opt	in	out	source	
							destination	
Chain FORWARD (policy DROP 0 packets, 0 bytes)								
pkts	bytes	target	prot	opt	in	out	source	
							destination	
5	0	ACCEPT	tcp	--	eth0	eth1	anywhere	192.168.60.
	0	ACCEPT	tcp	--	eth1	eth0	192.168.60.5	anywhere
	0	ACCEPT	tcp	--	eth1	eth1	192.168.60.5	anywhere
5	0	ACCEPT	tcp	--	eth1	eth1	anywhere	192.168.60.
	0	ACCEPT	tcp	--	eth1	eth1	anywhere	192.168.60.
6	0	ACCEPT	tcp	--	eth1	eth1	192.168.60.6	anywhere
	0	ACCEPT	tcp	--	eth1	eth1	192.168.60.7	anywhere
7	0	ACCEPT	tcp	--	eth1	eth1	anywhere	192.168.60.
	0	ACCEPT	tcp	--	eth1	eth1	anywhere	192.168.60.
Chain OUTPUT (policy DROP 0 packets, 0 bytes)								
pkts	bytes	target	prot	opt	in	out	source	
							destination	
root@bf0a51fb96a8:/# █								

Upload screenshots showing proof of checking rules 1-5 above (Please label each items)

▼ 2C\_192.168.60.5\_telnet\_outside\_and\_internal.png

 Download

```
[05/03/24]seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcde host1-192.168.60.5
[05/03/24]seed@VM:~/.../Labsetup$ docksh 2c
root@2c5d5fccbcde:/# telnet 10.9.0.5
Trying 10.9.0.5...
^C
root@2c5d5fccbcde:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
267760e3bc58 login: skdmf
Password:
Connection closed by foreign host.
root@2c5d5fccbcde:/# telnet 192.168.60.7
Trying 192.168.60.7...
Connected to 192.168.60.7.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
547ae6becaad login: dklfj
Password:
^CConnection closed by foreign host.
root@2c5d5fccbcde:/#
```

▼ 2C\_192.168.60.6\_telnet\_outside\_and\_internal.png

 Download

```
[05/03/24] seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3  hostA-10.9.0.5
267760e3bc58  host2-192.168.60.6
547ae6becaad  host3-192.168.60.7
bf0a51fb96a8  seed-router
2c5d5fccbcfdc host1-192.168.60.5
[05/03/24] seed@VM:~/.../Labsetup$ docksh 26
root@267760e3bc58:/# telnet 10.9.0.5
Trying 10.9.0.5...
^C
root@267760e3bc58:/# telnet 192.168.60.7
Trying 192.168.60.7...
Connected to 192.168.60.7.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
547ae6becaad login: exit
Password:
^CConnection closed by foreign host.
root@267760e3bc58:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
2c5d5fccbcfdc login: cc
Password:
^CConnection closed by foreign host.
root@267760e3bc58:/# █
```

```
[05/03/24]seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcdc host1-192.168.60.5
[05/03/24]seed@VM:~/.../Labsetup$ docksh 54
root@547ae6becaad:# telnet 10.9.0.5
Trying 10.9.0.5...
^C
root@547ae6becaad:# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
267760e3bc58 login: dkds
Password:
^CConnection closed by foreign host.
root@547ae6becaad:# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
2c5d5fccbcdc login: sdawd
Password:
^CConnection closed by foreign host.
root@547ae6becaad:# █
```

▼ 2C\_outside\_telnet\_internal.png

 Download

```
[05/03/24] seed@VM:~/.../Labsetup$ dockps
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcde host1-192.168.60.5
[05/03/24] seed@VM:~/.../Labsetup$ docksh 9a
root@9aa7af02f8b3:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
^C^]
Ubuntu 20.04.1 LTS
^]
2c5d5fccbcde login: ^^^^^^^exit
Password:
^CConnection closed by foreign host.
root@9aa7af02f8b3:/# telnet 192.168.60.7
Trying 192.168.60.7...
^C
root@9aa7af02f8b3:/# telnet 192.168.60.6
Trying 192.168.60.6...
^C
root@9aa7af02f8b3:/#
```

When you are done with this task, please remember to clean the table (e.g. `iptables -F`) or restart the container before moving on to the next task.

### **Q3 Task 3: Connection Tracking and Stateful Firewall**

**25 Points**

In the previous task, we have only set up stateless firewalls, which inspect each packet independently. However, packets are usually not independent; they may be part of a TCP connection, or they may be ICMP packets triggered by other packets. Treating them independently does not take into consideration the context of the packets, and can thus lead to inaccurate, unsafe, or complicated firewall rules. For example, if we would like to allow TCP packets to get into our network only if a connection was made first, we cannot achieve that easily using stateless packet filters, because when the firewall examines each individual TCP packet, it has no idea whether the packet belongs to an existing connection or not, unless the firewall maintains some state information for each connection. If it does that, it becomes a stateful firewall.

### Q3.1 Task 3.A: Experiment with the Connection Tracking

15 Points

To support stateful firewalls, we need to be able to track connections. This is achieved by the conntrack mechanism inside the kernel. In this task, we will conduct experiments related to this module, and get familiar with the connection tracking mechanism. In our experiment, we will check the connection tracking information on the router container. This can be done using the following command:

```
# conntrack -L
```

The goal of the task is to use a series of experiments to help students understand the connection concept in this tracking mechanism, especially for the ICMP and UDP protocols, because unlike TCP, they do not have connections. Please conduct the following experiments. For each experiment, please describe your observation, along with your explanation.

**Hint:** Use `sysctl net.netfilter` to see the timeout settings, and verify.

#### Expected Output for ICMP only (similar output for UDP and TCP)

```
root@9193e8603cb9:/# conntrack -L
icmp      1 7 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=71 src=192.168.60.5
dst=10.9.0.5 type=0 code=0 id=71 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@9193e8603cb9:/# conntrack -L
icmp      1 0 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=71 src=192.168.60.5
dst=10.9.0.5 type=0 code=0 id=71 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@9193e8603cb9:/# conntrack -L
conntrack v1.4.5 (conntrack-tools): 0 flow entries have been shown.
root@9193e8603cb9:/#
```

- ICMP experiment: Run the following command and check the connection tracking information on the router. Describe your observation. How long is the ICMP connection state be kept?

```
// On 10.9.0.5, send out ICMP packets
# ping 192.168.60.5
```

**Upload your screenshots for ICMP connection tracking**

▼ 3A\_connection\_tracking\_icmp.png

 Download

```
root@bf0a51fb96a8:/# conntrack -L
icmp      1 27 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=119 src=192.168.60
.5 dst=10.9.0.5 type=0 code=0 id=119 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
icmp      1 26 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=119 src=192.168.60
.5 dst=10.9.0.5 type=0 code=0 id=119 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
icmp      1 25 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=119 src=192.168.60
.5 dst=10.9.0.5 type=0 code=0 id=119 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
icmp      1 24 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=119 src=192.168.60
.5 dst=10.9.0.5 type=0 code=0 id=119 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
icmp      1 23 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=119 src=192.168.60
.5 dst=10.9.0.5 type=0 code=0 id=119 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
icmp      1 22 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=119 src=192.168.60
.5 dst=10.9.0.5 type=0 code=0 id=119 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

**Describe your observation for ICMP connection tracking.** Be sure to specify how long the tracking information is kept once a connection has been established.

The ICMP connection tracking information is kept for 30 seconds

- UDP experiment: Run the following command and check the connection tracking information on the router. Describe your observation. How long is the UDP connection state be kept?

```
// On 192.168.60.5, start a netcat UDP server
# nc -lu 9090
// On 10.9.0.5, send out UDP packets
# nc -u 192.168.60.5 9090
<type something, then hit return>
```

**Upload your screenshots for UDP connection tracking**

▼ 3A\_connection\_tracking\_udp.png

 Download

```
root@bf0a51fb96a8:/# conntrack -L
udp      17 25 src=10.9.0.5 dst=192.168.60.5 sport=43285 dport=9090 [UNREPLIED]
src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=43285 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
udp      17 23 src=10.9.0.5 dst=192.168.60.5 sport=43285 dport=9090 [UNREPLIED]
src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=43285 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
udp      17 22 src=10.9.0.5 dst=192.168.60.5 sport=43285 dport=9090 [UNREPLIED]
src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=43285 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
udp      17 21 src=10.9.0.5 dst=192.168.60.5 sport=43285 dport=9090 [UNREPLIED]
src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=43285 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
conntrack v1.4.5 (conntrack-tools): 0 flow entries have been shown.
root@bf0a51fb96a8:/# █
```

**Describe your observation for UDP connection tracking.** Be sure to specify how long the tracking information is kept once a connection has been established.

The UDP connection tracking information is kept for 30 seconds

- TCP experiment: Run the following command and check the connection tracking information on the router. Describe your observation. Be sure to specify how long the tracking information is kept once a connection has been established.

```
// On 192.168.60.5, start a netcat TCP server
# nc -l 9090
// On 10.9.0.5, send out TCP packets
# nc 192.168.60.5 9090
<type something, then hit return>
```

**Upload your screenshots for TCP connection tracking**

▼ 3A\_connection\_tracking\_tcp.png

 Download

```
root@bf0a51fb96a8:/# conntrack -L
tcp      6 431985 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=48806 dport=9
090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=48806 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# conntrack -L
tcp      6 431995 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=48806 dport=9
090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=48806 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@bf0a51fb96a8:/# █
```

**Describe your observation for TCP connection tracking.** Be sure to specify how long the tracking information is kept.

The UDP connection tracking information is kept for 432000 seconds

### **Q3.2 Task 3.B: Setting Up a Stateful Firewall**

**10 Points**

Now we are ready to set up firewall rules based on connections. In the following example, the "-m conntrack" option indicates that we are using the conntrack module, which is a very important module for iptables; it tracks connections, and iptables relies on the tracking information to build stateful firewalls. The --ctstate ESTABLISHED,RELATED indicates that whether a packet belongs to an ESTABLISHED or RELATED connection. The rule allows TCP packets belonging to an existing connection to pass through.

```
iptables -A FORWARD -p tcp -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
```

The rule above does not cover the SYN packets, which do not belong to any established connection. Without it, we will not be able to create a connection in the first place. Therefore, we need to add a rule to accept incoming SYN packet:

```
iptables -A FORWARD -p tcp -i eth0 --dport 8080 --syn -m conntrack --ctstate NEW -j ACCEPT
```

Finally, we will set the default policy on FORWARD to drop everything. This way, if a packet is not accepted by the two rules above, they will be dropped.

```
iptables -P FORWARD DROP
```

Please rewrite the firewall rules in Q2.3 Task 2.C, but this time, we will add a rule allowing internal hosts to visit any external server (this was not allowed in Task 2.C). After you write the rules using the connection tracking mechanism, think about how to do it without using the connection tracking mechanism (you do not need to actually implement them). When you are done with this task, remember to clear all the rules.

**Upload your IPTTables rules** `iptables -L -v` **and highlight the new rule to allow internal hosts to visit any external server**

▼ 3B\_rules.png

 Download

```
root@bf0a51fb96a8:/# iptables -L -v
Chain INPUT (policy DROP 0 packets, 0 bytes)
pkts bytes target prot opt in     out      source          destination
Chain FORWARD (policy DROP 6 packets, 374 bytes)
pkts bytes target prot opt in     out      source          destination
  140  7914 ACCEPT  tcp  --  any    any    anywhere       anywhere        ctstate RELATED,ESTABLISHED
      1   60 ACCEPT  tcp  --  eth0   eth1   anywhere       192.168.60.5      tcp dpt:telnet
      0   0 ACCEPT  tcp  --  eth1   eth1   anywhere       anywhere        ctstate RELATED,ESTABLISHED
      3   180 ACCEPT  tcp  --  eth1   eth0   anywhere       anywhere        ctstate NEW
      0   0 ACCEPT  tcp  --  eth0   eth1   anywhere       anywhere        ctstate NEW
Chain OUTPUT (policy DROP 0 packets, 0 bytes)
pkts bytes target prot opt in     out      source          destination
root@bf0a51fb96a8:/#
```

## Q4 Task 4: Limiting Network Traffic

10 Points

In addition to blocking packets, we can also limit the number of packets that can pass through the firewall. This can be done using the limit module of iptables. In this task, we will use this module to limit how many packets from 10.9.0.5 are allowed to get into the internal network. You can use "iptables -m limit -h" to see the manual.

```
$ iptables -m limit -h
limit match options:
--limit avg      max average match rate: default 3/hour
                  [Packets per second unless followed by
                   /sec /minute /hour /day postfixes]
--limit-burst number  number to match in a burst, default 5
```

Please run the following commands on router, and then ping 192.168.60.5 from 10.9.0.5. Describe your observation. Please conduct the experiment with and without the second rule, and then explain whether the second rule is needed or not, and why.

```
iptables -A FORWARD -s 10.9.0.5 -m limit \
          --limit 10/minute --limit-burst 5 -j ACCEPT
iptables -A FORWARD -s 10.9.0.5 -j DROP
```

### Expected Output with second rule

```
root@33bd6f6ddcae:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.083 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.135 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.161 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.111 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.113 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.105 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.111 ms
64 bytes from 192.168.60.5: icmp_seq=19 ttl=63 time=0.116 ms
^C
--- 192.168.60.5 ping statistics ---
24 packets transmitted, 8 received, 66.6667% packet loss, time 2353
1ms
rtt min/avg/max/mdev = 0.083/0.116/0.161/0.021 ms
```

Upload your screenshots below

```
root@9aa7af02f8b3:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.1
16 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.0
55 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.1
03 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.2
56 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.1
03 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.1
50 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.
155 ms
64 bytes from 192.168.60.5: icmp_seq=19 ttl=63 time=0.
056 ms
```

**Explain your observation on whether the second rule is needed or not, and why**

After the initial burst of 5 packets, the firewall starts dropping packets so that the limit of 10 packets per minute is upheld. The second rule is needed because otherwise the firewall does not drop packets so that the limit is upheld (since the default rule on FORWARD is to accept everything).

## **Q5 Task A - Multiple Choice**

**5 Points**

**Q5.1**

**1 Point**

This Chain applies for all packets that are addressed to the firewall

- INPUT
- OUTPUT
- FORWARD

**Q5.2**

**1 Point**

This chain applies for all packets originating from firewall and going out of the server

- INPUT
- OUTPUT
- FORWARD

**Q5.3**

**1 Point**

This chain applies for all packets passing through the firewall from other hosts on the network. The host with iptables is neither the source nor destination of the packet; mainly used to route packets through the machines on the network.

- INPUT
- OUTPUT
- FORWARD

**Q5.4****1 Point**

This jump target does not respond to a packet at all and does nothing with the packet. If an attack sends a packet, they would not get any response.

DROP

REJECT

**Q5.5****1 Point**

This jump target responds with an ICMP Destination Unreachable back to the source. This indicates that a server exists, which is beneficial for troubleshooting and for attackers.

DROP

REJECT

## Q6 Task B (not in SEED Lab PDF)

5 Points

### iptables Logging

**Note:** this exercise is performed on the Guest VM (the first VM that you log into, not the docker containers. Your prompt should say `seed@[hostname]`). Also, there are other rules here, ignore them, or if they get messed up, restart the VM to restore the rules.

Sometimes you need to log certain packets. Using the jump target LOG (`-j LOG`). Write a rule to log all outgoing icmp packets sent to 8.8.8.8 with the message "ICMP to Google!" (using `--log-prefix`)

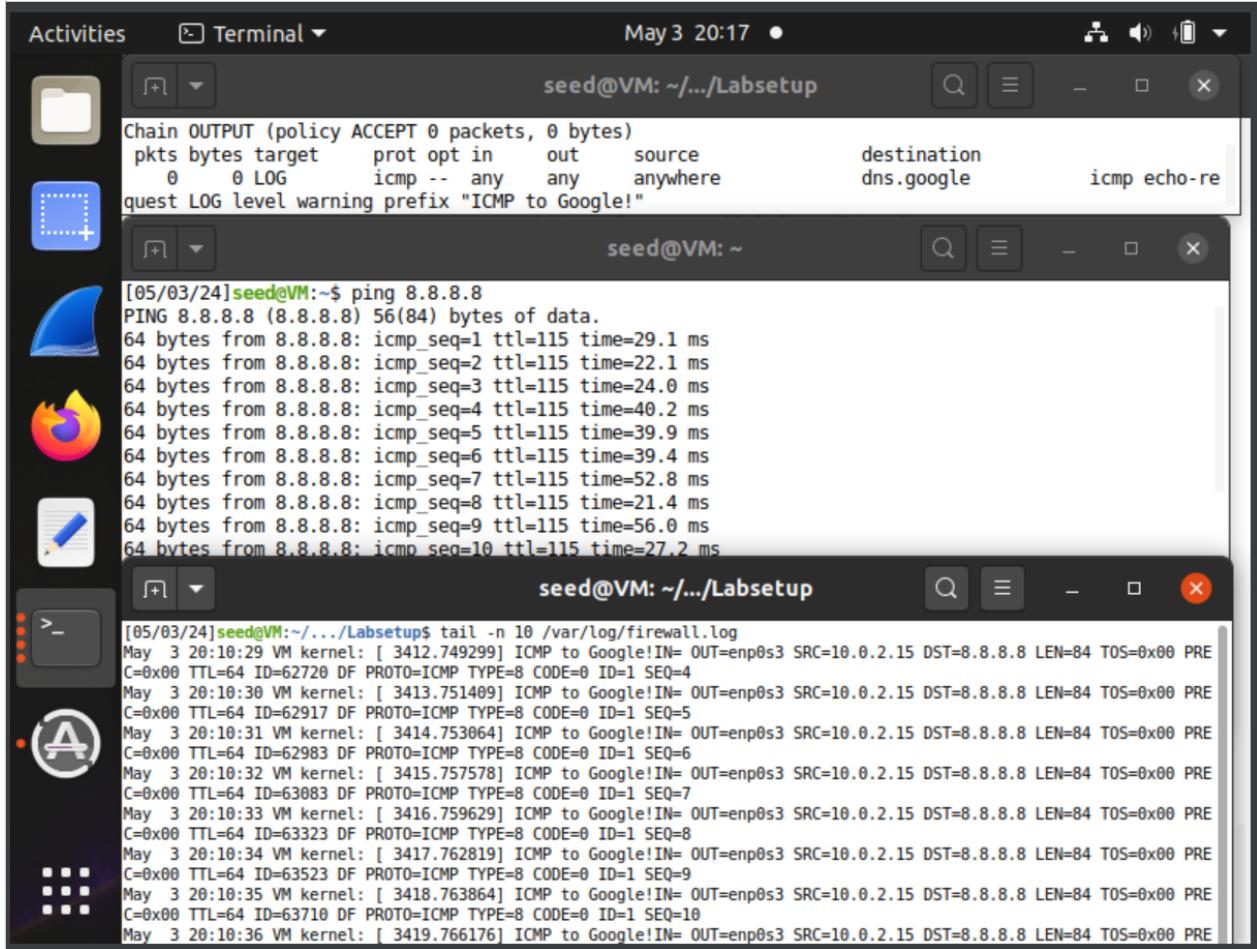
1. Enable logging by adding the line `kern.warn /var/log/firewall.log` to the file `/etc/rsyslog.conf` (be sure to use sudo)
2. Restart rsyslog: `sudo service rsyslog restart`
3. Create traffic to be logged (using an iptables rule with `-j LOG`)
4. The log is saved to `/var/log/firewall.log`

**Submit a screenshot with the following** (full VM with date and time):

1. Your `iptables -L -v` with the LOG rule
2. Your ping from the terminal
3. Log file showing the pings (you can use something such as  
`tail -n 10 /var/log/firewall.log`)

▼ Task B.png

 Download



```
Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
pkts bytes target     prot opt in     out     source               destination
      0    0 LOG       icmp --  any    any    anywhere             dns.google          icmp echo-request
quest LOG level warning prefix "ICMP to Google!"
```

```
[05/03/24]seed@VM:~$ ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=115 time=29.1 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=115 time=22.1 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=115 time=24.0 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=115 time=40.2 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=115 time=39.9 ms
64 bytes from 8.8.8.8: icmp_seq=6 ttl=115 time=39.4 ms
64 bytes from 8.8.8.8: icmp_seq=7 ttl=115 time=52.8 ms
64 bytes from 8.8.8.8: icmp_seq=8 ttl=115 time=21.4 ms
64 bytes from 8.8.8.8: icmp_seq=9 ttl=115 time=56.0 ms
64 bytes from 8.8.8.8: icmp_seq=10 ttl=115 time=27.2 ms
```

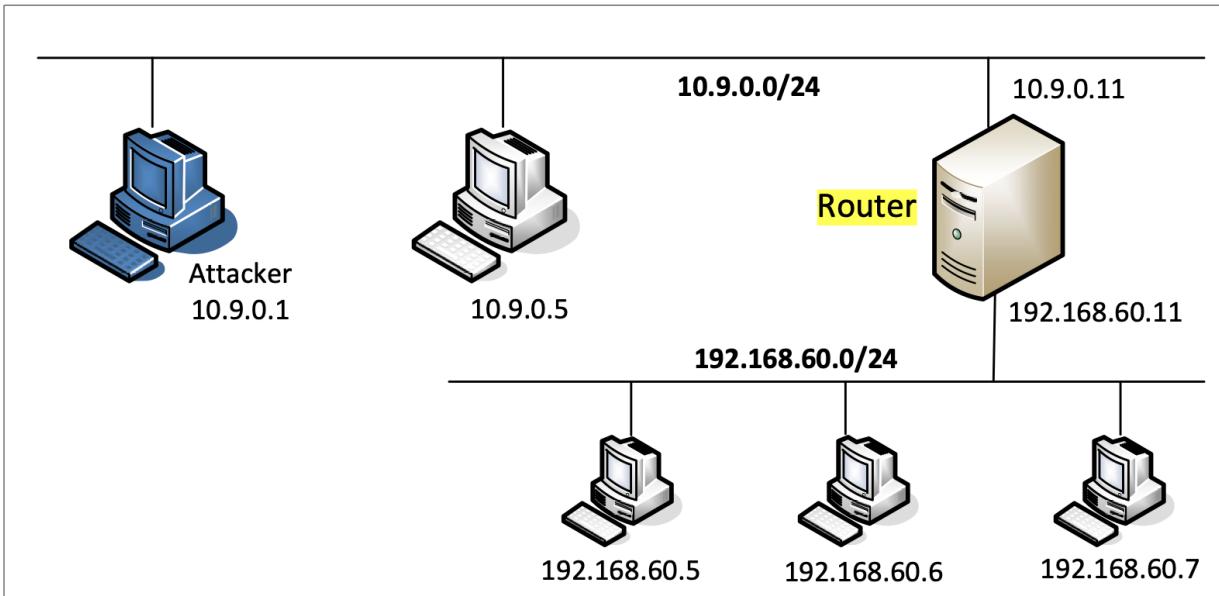
```
[05/03/24]seed@VM:~/Labsetup$ tail -n 10 /var/log/firewall.log
May  3 20:10:29 VM kernel: [ 3412.749299] ICMP to Google!IN= OUT=enp0s3 SRC=10.0.2.15 DST=8.8.8.8 LEN=84 TOS=0x00 PREC=0x00 TTL=64 ID=62720 DF PROTO=ICMP TYPE=8 CODE=0 ID=1 SEQ=4
May  3 20:10:30 VM kernel: [ 3413.751409] ICMP to Google!IN= OUT=enp0s3 SRC=10.0.2.15 DST=8.8.8.8 LEN=84 TOS=0x00 PREC=0x00 TTL=64 ID=62917 DF PROTO=ICMP TYPE=8 CODE=0 ID=1 SEQ=5
May  3 20:10:31 VM kernel: [ 3414.753064] ICMP to Google!IN= OUT=enp0s3 SRC=10.0.2.15 DST=8.8.8.8 LEN=84 TOS=0x00 PREC=0x00 TTL=64 ID=62983 DF PROTO=ICMP TYPE=8 CODE=0 ID=1 SEQ=6
May  3 20:10:32 VM kernel: [ 3415.757578] ICMP to Google!IN= OUT=enp0s3 SRC=10.0.2.15 DST=8.8.8.8 LEN=84 TOS=0x00 PREC=0x00 TTL=64 ID=63083 DF PROTO=ICMP TYPE=8 CODE=0 ID=1 SEQ=7
May  3 20:10:33 VM kernel: [ 3416.759629] ICMP to Google!IN= OUT=enp0s3 SRC=10.0.2.15 DST=8.8.8.8 LEN=84 TOS=0x00 PREC=0x00 TTL=64 ID=63323 DF PROTO=ICMP TYPE=8 CODE=0 ID=1 SEQ=8
May  3 20:10:34 VM kernel: [ 3417.762819] ICMP to Google!IN= OUT=enp0s3 SRC=10.0.2.15 DST=8.8.8.8 LEN=84 TOS=0x00 PREC=0x00 TTL=64 ID=63523 DF PROTO=ICMP TYPE=8 CODE=0 ID=1 SEQ=9
May  3 20:10:35 VM kernel: [ 3418.763864] ICMP to Google!IN= OUT=enp0s3 SRC=10.0.2.15 DST=8.8.8.8 LEN=84 TOS=0x00 PREC=0x00 TTL=64 ID=63710 DF PROTO=ICMP TYPE=8 CODE=0 ID=1 SEQ=10
May  3 20:10:36 VM kernel: [ 3419.766176] ICMP to Google!IN= OUT=enp0s3 SRC=10.0.2.15 DST=8.8.8.8 LEN=84 TOS=0x00 PREC=0x00 TTL=64 ID=63907 DF PROTO=ICMP TYPE=8 CODE=0 ID=1 SEQ=11
```

When done, remove the rule ([-D](#)), or just restart the entire VM.

## Q7 Task C (not in the SEED Lab PDF)

20 Points

Flush all your rules (`iptables -F`), and write the following rules **only** on the Router.



All traffic should be DROP'ed except for the following rules:

1. Any outgoing traffic from the **192.168.60.0/24** network to the **10.9.0.0/24** network is allowed. That is, the connection must be initiated from the **192.168.60.0/24** network (and not the other way around).
2. Any host can ping the router (**10.9.0.11** or **192.168.60.11**)
3. Host **10.9.0.5** is allowed to telnet to **192.168.60.5**

**All rules must be stateful**

Your complete iptable rules. **Please write the rules in the same order given above (1,2, then 3).** Hint: do not forget the default policy.

```
-----  
iptables -A FORWARD -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
```

```
1. iptables -A FORWARD -i eth1 -o eth0 -d 10.9.0.0/24 -s 192.168.60.0/24 -m  
conntrack --ctstate NEW -j ACCEPT
```

```
2. iptables -A INPUT -p icmp --icmp-type echo-request -m conntrack --ctstate  
NEW,ESTABLISHED -j ACCEPT
```

```
iptables -A OUTPUT -p icmp --icmp-type echo-reply -m conntrack --ctstate  
ESTABLISHED,RELATED -j ACCEPT
```

```
3. iptables -A FORWARD -p tcp -i eth0 -o eth1 -s 10.9.0.5 -d 192.168.60.5 --dport 23 --syn -m conntrack --ctstate NEW -j ACCEPT
```

```
iptables -P FORWARD DROP
```

```
iptables -P OUTPUT DROP
```

```
iptables -P INPUT DROP
```

Take a screenshot of `iptables -L -v`. Ensure your screenshot is of the entire VM.

▼ Task C.png  [Download](#)

```
9aa7af02f8b3 hostA-10.9.0.5
267760e3bc58 host2-192.168.60.6
547ae6becaad host3-192.168.60.7
bf0a51fb96a8 seed-router
2c5d5fccbcda host1-192.168.60.5
[05/04/24]seed@VM:~/.../Labsetup$ docksh bf
root@bf0a51fb96a8:/# iptables -L -v
Chain INPUT (policy DROP 0 packets, 0 bytes)
pkts bytes target     prot opt in     out     source               destination
      5   420 ACCEPT     icmp  --  any    any     anywhere             anywhere
          icmp echo-request ctstate NEW,ESTABLISHED

Chain FORWARD (policy DROP 4 packets, 254 bytes)
pkts bytes target     prot opt in     out     source               destination
      8   448 ACCEPT     all   --  any    any     anywhere             anywhere
          ctstate RELATED,ESTABLISHED
      0     0 ACCEPT     all   --  eth1   eth0   192.168.60.0/24    10.9.0.0/24
          ctstate NEW
      1     60 ACCEPT    tcp  --  eth0   eth1   10.9.0.5            192.168.60.
      5     5 ACCEPT    tcp  dpt:telnet flags:FIN,SYN,RST,ACK/SYN ctstate NEW

Chain OUTPUT (policy DROP 0 packets, 0 bytes)
pkts bytes target     prot opt in     out     source               destination
      5   420 ACCEPT     icmp  --  any    any     anywhere             anywhere
          icmp echo-reply ctstate RELATED,ESTABLISHED
root@bf0a51fb96a8:/#
```

## Q8 Early/Date Submission Bonus

0 Points

Bonus points for early or late submission will be added here. You may submit up to five days early for an extra 5% bonus points added to the grade of this assignment, or up to 4% deducted for late submission.

Submissions more than 4 days late are not accepted without an approved reason.